

CS211 Computer Architecture II

ISAs and Microprogramming Homework #0

Assigned 13/09/2022

Due 23:59:59 27/09/2022

<https://toast-lab.sist.shanghaitech.edu.cn/courses/CS211@ShanghaiTech/Fall-2022/>

The homework is intended to help you learn the material, and we encourage you to collaborate with other students and to ask questions in discussion sections and office hours to understand the problems. However, each student must turn in their own solution to the problems.

The homework also provides essential background material for the mid-term and final exams. It will be graded primarily on an effort basis, but if you do not work through it, you are unlikely to succeed on the mid-term and final exams! We will distribute solutions to homework assignment. Note that each homework assignment is due at its respective due time, and all assignments are to be submitted **in English**. However, late submissions will **NOT** be accepted, except for extreme circumstances and/or with prior arrangement.

Name: Zhenghong Yu 俞政宏
ID: 2020533156

Problem 1: CISC, RISC, and Stack: Comparing ISAs (3 points)

In this problem, your task is to compare three different ISAs: x86 (a CISC architecture with variable-length instructions), RISC-V (a load-store, RISC architecture with 32-bit instructions in its base form), and a stack-based ISA.

Problem 1.A

CISC (1 point)

Let us begin by considering the following C code:

```
int bar (int x, int y, int z) {
    for (; x >= y; --x)
        z += x;
    return z;
}
```

X=15 y=3 z=1

Using gcc and objdump on an x86 machine, we see that the above loop compiles to the following x86 instruction sequence. On entry to this code, register %eax, %ebx and %ecx contains x, y, and z respectively. Throughout parts (1.A--1.C), we will ignore what happens in the done label and return statement.

```
start:
    cmp    %ebx, %eax   6 bytes
    jl     done          2 bytes
loop:
    add    %eax, %ecx   2 bytes
    dec    %eax          2 bytes
    cmp    %eax, %ebx   6 bytes
    jle    loop          2 bytes
done:
    ret               1 bytes
```

The meanings and instruction lengths of the instructions used above are given in the following table. Registers are denoted with $R_{\text{SUBSCRIPT}}$, register contents with $\langle R_{\text{SUBSCRIPT}} \rangle$.

Instruction	Operation	Length
add $R_{\text{SRC}}, R_{\text{DEST}}$	$\langle R_{\text{DEST}} \rangle + \langle R_{\text{SRC}} \rangle \rightarrow R_{\text{DEST}}$	2 bytes
dec R_{DEST}	$R_{\text{DEST}} - 1 \rightarrow R_{\text{DEST}}$	2 bytes
cmp $R_{\text{SRC1}}, R_{\text{SRC2}}$	$\langle R_{\text{SRC2}} \rangle - \langle R_{\text{SRC1}} \rangle \rightarrow \text{Temp}$	6 bytes
jl label	if (SF != OF) jump to the address specified by label	2 bytes
jle label	if ((SF != OF) or (ZF == 1)) jump to the address specified by label	2 bytes
ret	return to the caller function	1 byte

Notice that the jump instructions `j l` and `j l e` (jump if less, and jump if less or equal) depend on status flags, i.e., SF, ZF and OF in the table. Status flags are set by the instruction preceding a jump, based on the result of the computation. The meanings of the status flags are given in the following table:

Name	Purpose	Condition Reported
ZF	Zero flag	Result is zero (<code>cmp</code> performs an actual subtraction)
SF	Sign flag	Set equal to high-order bit of result (0 if positive 1 if negative)
OF	Overflow flag	Set if result is too large a positive number or too small a negative number (excluding sign bit) to fit in destination operand; cleared otherwise.

How many bytes is the program? For the above x86 assembly code, how many bytes of instructions need to be fetched if $x = 13$, $y = 3$, and $z = 1$? Assuming 32-bit data values, how many bytes of data memory need to be loaded? Stored?

- (1) 21 bytes
- (2) 141 bytes
- (3) load: 0 bytes
store: 0 bytes

Problem 1.B

RISC (1 points)

Translate each of the x86 instructions in the following table into corresponding RISC-V instructions. Place the `loop` label where appropriate. You should use the minimum number of instructions needed to translate each x86 instruction. Assume that upon entry, $x1$ contains x , $x2$

contains y, and x3 contains z. If needed, use x4 as a condition register, and x5, x6, etc., for temporaries. You should not need to use any floating-point registers or instructions in your code. A description of the RISC-V instruction set architecture can be found in the class website alongside L03.

x86 instruction	label	RISC-V instruction sequence	
	start:		
cmp %ebx, %eax		blt x1, x2, done	4
j1 done	loop:		
add %eax, %ecx		add x3, x1, x3	4
dec %eax		addi x1, x1, -1	4
cmp %eax, %ebx		bge x1, x2, loop	4
jle loop			
	done:		
ret		ret	4

How many bytes is the RISC-V program using your direct translation? How many bytes of RISC-V instructions need to be fetched for x = 13, y = 3, and z = 1 using your direct translation? Assuming 32-bit data values, how many bytes of data memory need to be loaded? Stored?

- (1) 20 bytes
- (2) 140 bytes
- (3) load: 0 bytes
Store: 0 bytes

Problem 1.C**Stack (1 points)**

In a stack architecture, all operations occur on top of the stack. Only push and pop access memory, and all other instructions remove their operands from the stack and replace them with the result. The hardware implementation we will assume for this problem uses stack registers for the top two entries; accesses that involve other stack positions (e.g., pushing or popping something when the stack has more than two entries) use an extra memory reference. Assume each instruction occupies three bytes if it takes an address or label; other instructions occupy one byte.

Instruction	Definition
PUSH A	load value at M[A]; push value onto stack 3
POP A	pop stack; store value to M[A] 3
ADD	pop two values from the stack; ADD them; push result onto stack
SUB	pop two values from the stack; SUBtract top value from the 2nd; push result onto stack
INC	pop value from top of stack; increment value by one; push result onto stack
ZERO	zero out value at top of stack
DEC	pop value from top of stack; decrement value by one; push result onto stack
BEQZ <i>label</i>	pop value from stack; if it's zero, continue at <i>label</i> ; else, continue with next instruction
BNEZ <i>label</i>	pop value from stack; if it's not zero, continue at <i>label</i> ; else, continue with next instruction
BNEG <i>label</i>	pop value from stack; if it's negative (less than zero), continue at <i>label</i> ; else, continue with next instruction
JUMP <i>label</i>	continue execution at location <i>label</i>

Translate the `bar` loop to the stack ISA. For uniformity, please use the same control flow as in parts (1.A) and (1.B). Assume that when we reach the loop, `x`, `y`, and `z` are at the top of the stack. At the end of the loop, `z` should be at the top of the stack.

How many bytes is the stack program using your translation? How many bytes of instructions need to be fetched for `x = 13`, `y = 3` and `z = 1` using your translation? Assuming 32-bit data values, how many bytes of data memory need to be loaded? Stored? Would the number of bytes loaded and stored change if the stack could fit 8 entries in registers?

x
y
z
l

Start: POP X 3
 POP Y 3
 PUSH Y 3
 PUSH X 3
 SUB 1

(1) 38 bytes
 (2) 258 bytes
 (3) load: 184 bytes
 store: 52 bytes

16 BNEG done 3

(4) Don't change

loop: PUSH X 3
 ADD 1
 PUSH X 3
 DEC 1
 POP X 3
 PUSH Y 3
 PUSH X 3
 SUB 1
 DEC 1

22

BNEG loop 3

done: