

Advanced Computer Architecture: Lab 0 Report

Due on September 23rd at 11:59am

Name: **Zhenghong Yu**
Mailbox: **yuzhh1@shanghaitech.edu.cn**
Student ID: 2020533156

Contents

Realize LR and SC in simulator	3
Requirement of LR and SC	3
Premise	3
Details	3
Idea	3
Implementation	4
Realize CAS in lab0.c	10
Requirement of CAS	10
Details	10
Idea	10
Implementation	10
Result	12
Verification	12

Realize LR and SC in simulator

Requirement of LR and SC

Premise

The content below based on the "A" **Standard Extension for Atomic Instructions, Version 2.1** which is the newest version on the wiki.riscv.org.

Details

- LR and SC are both 32-bits instructions with opcode `0x2F`, funct5 (which means the last 5 bits number of the instruction) is `0x3` for 8 bytes data and `0x2` for 4 bytes data. Additionally, we don't need to consider aligned data, number limitation between LR and SC.

- LR loads a data from the address in `rs1`, places the **sign-extended** value in `rd`, and registers a **reservation set**—a set of bytes that subsumes the bytes in the addressed word.

- SC conditionally writes a data in `rs2` to the address in `rs1`: the SC succeeds only if the reservation is still **valid** and the reservation set contains the bytes being written. If the SC succeeds, the instruction writes the data in `rs2` to memory, and it writes **zero** to `rd`. If the SC fails, the instruction does **not** write to memory, and it writes a **nonzero** value to `rd`. Regardless of success or failure, executing an SC instruction invalidates **any** reservation held by this hart.

- An SC can only **pair with** the **most recent** LR in program order. An SC may succeed only if **no store** from another hart to the reservation set can be observed to have occurred between the LR and the SC, and if there is **no other** SC between the LR and itself in program order. An SC may succeed only if no write from a device other than a hart to the bytes accessed by the LR instruction can be observed to have occurred between the LR and SC. Note this LR might have had a different effective address and data size, but reserved the SC's address as part of the reservation set.

- The SC must **fail** if the address is not within the reservation set of the **most recent** LR in program order. The SC must **fail** if a store to the reservation set from another hart can be observed to occur between the LR and SC. The SC must **fail** if a write from some other device to the bytes accessed by the LR can be observed to occur between the LR and SC. An SC must **fail** if there is another SC (to any address) between the LR and the SC in program order.

- The failure code with value **1** is reserved to encode an unspecified failure. Other failure codes are reserved at this time, and portable software should only assume the failure code will be non-zero.

- To other access memory operations (**Load** or **Store**), we only concern about store instruction which access to these addresses that have been already reserved, remove these addresses from the reservation set.

Idea

- I use **registry_addr** to record the first address, **registry_whe** to record whether the reservation set is still valid, **registry_num** to record the length of the reservation set. These are declared in the *Simulator.h*.

- I almost use the datapath of **Load** instruction and **Store** instruction, as they have the same hazard

which can handle samely.

· Specially, despite the instruction fetch and datapath set, I handle the most important part (reserve and check) in the *classmemoryAccess*, which I think it might be more easier to be realize on hardware.

Implementation

New variable or modification statement in *Simulator.h*

```

1      .....
2      namespace RISCV
3      {
4          .....
5          enum Inst
6          {
7              .....
8              LRW = 54,    /* add encode lr.w instruction as 54 */
9              LRD = 55,    /* add encode lr.d instruction as 55 */
10             SCW = 56,    /* add encode sc.w instruction as 56 */
11             SCD = 57,    /* add encode sc.d instruction as 57 */
12             .....
13         };
14         .....
15         /* add LR SC instruction with opcode 0X2F */
16         const int OP_CAS = 0x2F;
17         .....
18         inline bool isReadMem(Inst inst)
19         {
20             /* As LR and SC both need memory access and write back to
21              reg, so we need to handle with harzard in pipline */
22             if (..... || inst == LRW || inst == LRD
23                 || inst == SCW || inst == SCD)
24             {
25                 return true;
26             }
27             return false;
28         }
29     }
30
31     class Simulator
32     {
33     public:
34         .....
35     private:
36         .....
37         struct EReg
38         {
39             .....

```

```

40         bool _cas;
41     } eReg, eRegNew;
42     .....
43     uint32_t registry_addr; /* Registry hold the first address */
44     uint32_t registry_num; /* Registry hold the number */
45     bool registry_whe; /* Registry hold the valid bit */
46 }

```

New variable or modification statement in *Simulator.cpp*

```

1     .....
2     namespace RISCV
3     {
4         .....
5         const char *INSTNAME[] { /* add instruction name */
6             ....., "lrw", "lrd", "scw", "scd"};
7     }
8     .....
9     void Simulator::decode()
10    {
11        .....
12        if (this->fReg.len == 4) // 32 bit instruction
13        {
14            .....
15            switch (opcode)
16            {
17                .....
18                case OP_CAS: /* add OP_CAS handler */
19                {
20                    /* read data in reg[rs1] which store the address */
21                    op1 = this->reg[rs1];
22                    /* read data in reg[rs2] which is no mean in LR,
23                     store the old_value in SC */
24                    op2 = this->reg[rs2];
25                    reg1 = rs1;
26                    reg2 = rs2;
27                    dest = rd; /* The reg which should be write back */
28
29                    /* Get the funct5 */
30                    uint32_t temp = (inst >> 27) & 0x1F;
31                    switch (temp)
32                    {
33                        case 0x2: /* LR instruction */
34                            if (funct3 == 3) /* LR.D instruction */
35                            {
36                                instname = "lrd";
37                                insttype = LRD;
38                            }
39                            else if (funct3 == 2) /* LR.W instruction */
40                            {

```

```

41         instname = "lrw";
42         insttype = LRW;
43     }
44     else
45     {
46         this->panic("Unknown 32 bit funct3
47                     0x%x\n", funct3);
48     }
49     break;
50 case 0x3: /* SC instruction */
51     if (funct3 == 3) /* SC.D instruction */
52     {
53         instname = "scd";
54         insttype = SCD;
55     }
56     else if (funct3 == 2) /* SC.W instruction */
57     {
58         instname = "scw";
59         insttype = SCW;
60     }
61     else
62     {
63         this->panic("Unknown 32 bit funct3
64                     0x%x\n", funct3);
65     }
66     break;
67 default:
68     this->panic("Unknown 32 bit funct5
69                 0x%x\n", temp);
70 }
71 }
72 .....
73 }
74 }
75 .....
76 }
77
78 void Simulator::execute()
79 {
80     .....
81     bool _cas = false; /* To identify whether it is a LR or SC */
82
83     switch (inst)
84     {
85         .....
86         case LRW: /* Datapath of LR.W instruction */
87             readMem = true;
88             writeReg = true;
89             memLen = 4;

```

```

90         out = op1;
91         readSignExt = true;
92         _cas = true;
93         break;
94     case LRD: /* Datapath of LR.D instruction */
95         readMem = true;
96         writeReg = true;
97         memLen = 8;
98         out = op1;
99         readSignExt = true;
100        _cas = true;
101        break;
102    case SCW: /* Datapath of SC.W instruction */
103        writeReg = true;
104        writeMem = true;
105        memLen = 4;
106        out = op1;
107        _cas = true;
108        op2 = op2 & 0xFFFFFFFF;
109        break;
110    case SCD: /* Datapath of SC.D instruction */
111        writeReg = true;
112        writeMem = true;
113        memLen = 8;
114        out = op1;
115        op2 = op2 & 0xFFFFFFFF;
116        _cas = true;
117        break;
118        .....
119    }
120    .....
121    /* Pass the information to the next stage */
122    this->eRegNew._cas = _cas;
123 }
124
125 void Simulator::memoryAccess()
126 {
127     .....
128     bool _cas = this->eReg._cas;
129     .....
130     if (writeMem)
131     {
132         switch (memLen)
133         {
134             case 1:
135                 good = this->memory->setByte(out, op2, &cycles);
136                 if ((out >= this->registry_addr) && (out <=
137                     (this->registry_addr + this->registry_num)))
138                 { /* If there is any byte write to the reservation set */

```

```

139         this->registry_whe = false;
140     }
141     break;
142 case 2:
143     good = this->memory->setShort(out, op2, &cycles);
144     for (uint32_t i = 0; i < 2; i++)
145     {
146         if (((out + i) >= this->registry_addr) && ((out + i)
147             <= (this->registry_addr + this->registry_num)))
148         { /* Any byte write to the reservation set */
149             this->registry_whe = false;
150         }
151     }
152     break;
153 case 4:
154     if (_cas)
155     { /* If this is SC instruction */
156         if (this->registry_whe && (this->registry_addr
157             == out) && (this->registry_num == 4))
158         { /* Reservation set is valid
159             the first address same
160             the access length same */
161             good = this->memory->setInt(out, op2, &cycles);
162             out = 0;
163         }
164         else
165         { /* If failed return code 1 */
166             out = 1;
167         }
168         /* SC eliminate all reservation */
169         this->registry_whe = false;
170     }
171     else
172     { /* If not, do basic store operation */
173         for (uint32_t i = 0; i < 4; i++)
174         { /* For all bytes */
175             if (((out + i) >= this->registry_addr) &&
176                 ((out + i) <= (this->registry_addr +
177                     this->registry_num)))
178             { /* Any access to the bytes */
179                 this->registry_whe = false;
180             }
181         }
182         good = this->memory->setInt(out, op2, &cycles);
183     }
184     break;
185 case 8:
186     if (_cas)
187     {

```



```

188         if (this->registry_whe && (this->registry_addr
189             == out) && (this->registry_num == 8))
190         {   /* Reservation set is valid
191             the first address same
192             the access length same */
193             good = this->memory->setLong(out, op2, &cycles);
194             out = 0;
195         }
196         else
197         {
198             out = 1;
199         }
200         tthis->registry_whe = false;
201     }
202     else
203     {
204         for (uint32_t i = 0; i < 8; i++)
205         {
206             if (((out + i) >= this->registry_addr)
207                 && ((out + i) <= (this->registry_addr
208                     + this->registry_num)))
209             {
210                 this->registry_whe = false;
211             }
212         }
213         good = this->memory->setLong(out, op2, &cycles);
214     }
215 }
216 }
217 .....
218 if (readMem)
219 {   /* Handle with LR instruction */
220     switch (memLen)
221     {
222         .....
223         case 4:
224             if(readSignExt)
225             {
226                 if(_cas)
227                 {   /* LR.W instruction */
228                     this->registry_whe = true;   /* Valid */
229                     this->registry_addr = out;   /* Store */
230                     this->registry_num = 4;   /* Length */
231                 }
232                 out = (int64_t) this->memory->getInt(out, &cycles);
233             }
234             .....
235         case 8:
236             if(readSignExt)

```

```

237         {
238             if(_cas)
239             { /* LR.W instruction */
240                 this->registry_whe = true;
241                 this->registry_addr = out;
242                 this->registry_num = 8;
243             }
244             out = (int64_t) this->memory->getInt(out, &cycles);
245         }
246         .....
247     }
248 }
249 }
250 }
251 }

```

Realize CAS in lab0.c

Requirement of CAS

Details

- Input of function is three variables: address which we want to store, old value that we think it should be, new value that we want to store in the address.
- We first load the value from the address, compare to the old value, if they are same (means until now, this address is not modified), store the new value into it and return code 1(true), if they are not same (means it has already been modified), stop store and return code 0(false).
- To make sure the CAS operation would not been interrupted, we must use **LR** instruction to load and **SC** instruction to store, in order to make sure the content of the particular address would not be modified while doing CAS operation. If **SC** failed, just try again until success.

Idea

- I use inline assembler to make sure that it will use *LR* and *SC* instruction. Also use keyword `__volatile__` to make sure that the assembler will not optimize my code.

Implementation

```

1 int CAS(long *dest, long new_value, long old_value)
2 {
3     int ret = 0;
4
5     // TODO: write your code here
6     __asm__ __volatile__(
7         "retry: lr.w %0, %[output2]\n" /* Load value of addr */

```

```
8      "      bne    %0, %[input1], fail\n" /* Compare and check */
9      "      sc.w   %0, %[input2], %[output2]\n" /* Store new value */
10     "      bnez   %0, retry\n" /* Whether atom operation success */
11     "      li     %[output1], 1\n" /* Return code 1 */
12     "      j      success\n" /* Go on following */
13     "fail: li     %[output1], 0\n" /* Return code 0 */
14     "success: \n"
15     : [output1] "=&r"(ret), [output2] "+A"(*dest)
16     : [input1] "rJ"(old_value), [input2] "rJ"(new_value)
17     : "memory");
18
19     return ret;
20 }
```

Result

I compile *lab0.c* and run it, the following result is correct (some debug output, just ignore it), the first circulation should be failed, the second circulation should be success and store new value into the address.

```
ubuntu@VM-4-7-ubuntu:~/RISCV-Simulator/build$ ./Simulator ../riscv-elf/lab0.riscv
Dst address is: 71608
CAS FAIL
0
1
Dst address is: 71608
CAS SUCCESS
1
211
Program exit from an exit() system call
----- STATISTICS -----
Number of Instructions: 418
Number of Cycles: 625
Avg Cycles per Instruction: 1.4952
Branch Prediction Accuracy: 0.5000 (Strategy: Always Not Taken)
Number of Control Hazards: 65
Number of Data Hazards: 209
Number of Memory Hazards: 20
-----
ubuntu@VM-4-7-ubuntu:~/RISCV-Simulator/build$
```

Verification

I check the whole process and write some simple other test to my implementation, I believe it is right.